

Chapel Compiler Overview

Steve Deitz

June, 2010

An important caveat. This document attempts to explain how the Chapel compiler works. It is most certainly incomplete and wrong.

Contents

1	Introduction	3
2	Intermediate Representation	3
2.1	BaseAST	3
2.2	Symbol	4
2.2.1	VarSymbol	5
2.2.2	ArgSymbol	5
2.2.3	TypeSymbol	5
2.2.4	FnSymbol	6
2.2.5	EnumSymbol	7
2.2.6	LabelSymbol	7
2.2.7	ModuleSymbol	7
2.3	Type	8
2.3.1	PrimitiveType	8
2.3.2	EnumType	9
2.3.3	ClassType	9
2.4	Expr	9
2.4.1	DefExpr	10
2.4.2	SymExpr	11
2.4.3	UnresolvedSymExpr	11
2.4.4	CallExpr	11
2.4.5	NamedExpr	11
2.4.6	BlockStmt	12
2.4.7	CondStmt	12
2.4.8	GotoStmt	12
2.5	Traversing the AST	13
2.5.1	The Core Traversal (With a Figure of the AST Graph)	13
2.5.2	The <code>getFirstExpr</code> and <code>getNextExpr</code> Methods	14
2.6	AST Functions and Methods	14
2.6.1	Methods for Insertion, Replacement, and Removal	15
2.6.2	Constructors and <code>new_Expr</code>	16
2.6.3	BaseAST Functions and Methods	17
2.6.4	Symbol Functions and Methods	18

2.6.5	Type Functions and Methods	18
2.6.6	Expr Functions and Methods	19
2.6.7	Stmt (Expr) Functions and Methods	19
2.6.8	Miscellaneous AST Utility Functions	20
2.6.9	Compilation Utility Routines	20
2.6.10	Optimizations and Analyses	21
2.7	Primitives	22
2.8	Pragmas	22
3	Passes	23
3.1	parse	23
3.2	checkParsed	23
3.3	cleanup	24
3.4	scopeResolve	24
3.5	flattenClasses	25
3.6	normalize	25
3.7	checkNormalized	26
3.8	buildDefaultFunctions	26
3.9	resolve	27
3.9.1	resolveFns (fn)	27
3.9.2	resolveBlock (fn)	28
3.9.3	resolveCall (call)	28
3.9.4	Wrappers	28
3.9.5	Generic Instantiation	28
3.9.6	Runtime Types	28
3.9.7	Scalar Promotion Types	29
3.9.8	Conditional Resolution and Try Tokens	29
3.9.9	Iterator Handling During Resolution	29
3.9.10	The Dream of Out-of-Order Resolution	30
3.9.11	The Dream of Combined Scope and Function Resolution	30
3.10	checkResolved	30
3.11	flattenFunctions	30
3.12	cullOverReferences	31
3.13	callDestructors	31
3.14	lowerIterators	31
3.14.1	Single Loop Iterator Optimization	32
3.14.2	Recursive Iterators	32
3.14.3	Iterators and Local Blocks	32
3.14.4	Miscellaneous Notes	33
3.15	parallel	33
3.16	prune	34
3.17	complex2record	34
3.18	removeUnnecessaryAutoCopyCalls	34
3.19	inlineFunctions	34
3.20	scalarReplace	34
3.21	refPropagation	34
3.22	copyPropagation	35

3.23	deadCodeElimination	35
3.24	removeWrapRecords	35
3.25	removeEmptyRecords	35
3.26	localizeGlobals	35
3.27	returnStarTuplesByRefArgs	35
3.28	insertWideReferences	35
3.29	optimizeOnClauses	36
3.30	insertLineNumbers	36
3.31	repositionDefExpressions	36
3.32	codegen	36
3.33	makeBinary	37
3.34	What Happens Between Passes?	37
4	Miscellaneous	37
4.1	Compiler Strings	37

1 Introduction

This document contains rough notes about the state of the compiler in June of 2010. It is subdivided into three sections. Section 2 describes the intermediate representation, sometimes referred to as the AST. Section 3 describes the passes and assumptions that may be made about the AST before and after the various passes. Section 4 covers miscellaneous topics.

This document assumes the reader has excellent familiarity with the Chapel language.

2 Intermediate Representation

The intermediate representation (IR), alternatively called the AST, is a graph-like structure defined by instances of subclasses of the `Symbol`, `Type`, and `Expr` classes (which are themselves subclasses of the `BaseAST` class). This representation is used throughout the entire compilation process, though assumptions on its structure change. (For example, after normalization, `CallExprs` will not be nested with the exception of a `CallExpr` that represents the `MOVE` primitive (C assignment).) The `BaseAST`, `Symbol`, `Type`, and `Expr` classes are never instantiated.

This graph is rooted at the module `ccrootModule` which contains `DefExprs` of all the modules in the program, as well as literals and other global entities.

Nodes in the representation can also be accessed via global vectors of each node type. These vectors are named as `gASTs` where `AST` is a particular AST node type. For example, `gFnSymbols` is a vector of all of the functions in a program, nested or not. Newly constructed nodes are automatically added to these vectors by their constructors. Between passes, these vectors are pruned of all AST nodes that are not part of the IR (part of the graph).

Three vectors of modules also capture the modules that make up the AST. The vector `allModules` is a vector of all the modules, and is probably identical to `gModuleSymbols`. The vector `userModules` contains all of the user modules. For debug purposes, this is useful for looking at the test code in question. The vector `mainModules` contains all of the main modules. See Section 2.2.7 for more information on modules.

2.1 BaseAST

Every node in the AST is a subclass of `BaseAST`. The `BaseAST` class contains three fields:

- `AstTag astTag` is an enumeration value used to identify a node's dynamic type. It is accessed via macros `isAST` and `toAST`, where `AST` is replaced by a node type, to check if a node is a particular type or to safely cast it to such a type (NULL if it is not such a type). This cast mechanism is used instead of C++'s `dynamic_cast` mechanism.

- `int id` is a unique integer assigned during construction. This is useful for debugging as each AST node will be numbered based on the order of construction. These numbers start at 1. During debugging with gdb via the `-gdb` flag, you can use the call `aid(id)` to get a pointer to an AST node based on this number.
- `int lineno` stores the line number of the Chapel code that this AST node originated from. During parsing, this field is set based on YACC state. Afterwards, these numbers are set based on the global variable `currentLineno`. This line number can be changed using the macro `SET_LINENO(ast)` which will set the global line number to the line number associated with a given AST node. Thus, for example, if you are adding a compiler-inserted `CallExpr` before an existing `CallExpr` in the AST, you can set the global line number to that of the existing `CallExpr`, and the new compiler-inserted AST nodes will use the same line number.

As a note for improvement, our line numbers are often incorrectly reported. If `SET_LINENO` were used more consistently, this would not be the case. One way to improve things would be to generate an internal error if `SET_LINENO` were not used correctly. To do this, we would want to unset the line number somehow. We could do this with an `UNSET_LINENO` macro or we could make the `SET_LINENO` macro require us to open a scope. Then when we leave this scope, we could set the current line number back to what it was before we opened this scope (and unset at the outermost level).

2.2 Symbol

The Symbol group of AST nodes derive from the class `Symbol`. Variables, constants, parameters, literals, arguments, fields, types, enumeration symbols, modules, and labels are all represented by symbols in the AST.

The fields in `Symbol` are:

- `const char* name` stores the Chapel name of the symbol. This is used during scope resolution and function resolution to resolve unresolved symbol uses (`UnresolvedSymExprs`). Comparing this field does not require a string comparison as it is canonicalized (see note on strings in Section 4.1).
- `const char* cname` stores the name that we will generate in C. This name is sometimes set during compilation to improve the readability of the generated code. During the codegen pass, we legalize these names and mangle them as necessary. Comparing this field does not require a string comparison as it is canonicalized (see note on strings in Section 4.1).
- `Type* type` stores the resolved type of variables, arguments, and fields. This is generally set during function resolution when types are resolved. Before such resolution, the type `dtUnknown` is used.
- `DefExpr* defPoint` points to a `DefExpr`, the expression hook in the AST. Every symbol has a `DefExpr` that connects the symbol to the AST. The symbol can also show up in the AST via the expression `SymExpr`. These account for both defs (writes to the symbols) and uses (reads of the symbols).
- `std::bitset<NUM_FLAGS> flags` is a bit vector of flags (referred to as pragmas when they appear in Chapel code). Flags are associated only with symbols. See Section 2.8 for more information about flags.

The subclasses of `Symbol` include the following:

- `VarSymbol` (Section 2.2.1)
- `ArgSymbol` (Section 2.2.2)
- `TypeSymbol` (Section 2.2.3)
- `FnSymbol` (Section 2.2.4)
- `EnumSymbol` (Section 2.2.5)
- `LabelSymbol` (Section 2.2.6)
- `ModuleSymbol` (Section 2.2.7)

2.2.1 VarSymbol

The VarSymbol class represents variables, including both global variables (in a module) and local variables (in a function). The VarSymbol is used for fields (in a type).

The constness of a variable is handled via a flag, so constants and parameters (compile-time constants) are represented by VarSymbol objects as well.

The fields of VarSymbol include the following:

- Immediate* `immediate` stores a representation of a literal symbol. For example, 1 is an integer literal. The representation of immediates is code that from IF1, an iterative flow analysis engine.

2.2.2 ArgSymbol

The ArgSymbol class is used to represent the formal arguments of functions. The fields of the ArgSymbol include the following:

- IntentTag `intent` stores the specified intent of the argument. This can be blank, in, inout, out, or const. Note that `INTENT_REF` does not appear to be used and probably predates the introduction of reference types. The intent field also marks generics via parameter and type intent arguments. The intent field is largely ignored after function resolution, which is to say that the functionality of the intent has been folded into the representation without it. Thus if the intent were inout, then a copy in and a copy out (via assignment and/or a move primitive) would have been added to the AST.
- BlockStmt* `typeExpr` stores the type of the argument as it is parsed in. This specified type is resolved during function resolution when the `type` field is set. After the `type` field is set, this field is ignored.
- BlockStmt* `defaultExpr` stores the expression specified as a default for the argument (what is used if no actual is specified). The default expression is folded into the AST during function resolution by creating a default wrapper function.
- BlockStmt* `variableExpr` stores the expression specified after an ellipsis specifying a variadic argument. This must resolve during function resolution to a parameter expression (it can also be omitted for variadic arguments or a query identifier). After function resolution, this field is ignored.
- Type* `instantiatedFrom` stores the argument from which this argument has been instantiated. That is, it points to the original argument in a generic function. It is only used during the function resolution pass.
- bool `instantiatedParam` marks an argument as a parameter that has been instantiated. This is only used during function resolution, and should most likely be a flag, or a hash table.
- bool `markedGeneric` marks an argument as generic (presumably). It is used to mark a type as generic if the Chapel programmer inserts a ? next to a generic type that has default values for all of its generic fields. This is also used to mark arguments when building default functions, but I'm not sure why. This should probably be a flag. It does span multiple passes (being set in `normalize` and used in `resolution`).

2.2.3 TypeSymbol

The TypeSymbol class is used for types. The subclasses of Type can be thought of as storing auxiliary information about types. There is a one-to-one mapping between TypeSymbols and Types. So the `type` field of TypeSymbol points to a subclass of Type, and the `symbol` field of this type class will point back at the TypeSymbol.

A good refactoring would probably be to use TypeSymbol everywhere that Type is used. Then a field in TypeSymbol could point to its auxiliary information, but we would never have to point back. This refactoring would require changing, for example, the `type` field of Symbol from Type to TypeSymbol.

The TypeSymbol class has no fields, inheriting all of the information it needs from the Symbol class, and putting all of the extra information in the Type class.

2.2.4 FnSymbol

The FnSymbol class is used to represent all methods and functions in a program. Top-level code in ModuleSymbols is moved into FnSymbols (their initialize functions pointed to by initFn in ModuleSymbol) during parsing.

The fields of FnSymbol include the following:

- `AList formals` is an alist (linked list of Expr subclasses) of DefExprs that define ArgSymbols. These are the formal arguments stored in a specified order.
- `DefExpr* setter` is non-null for var functions, and represents the setter argument. This argument is of bool type and is set to true only when the function is used in an lvalue setting. The implementation of var functions involves cloning the function so that there are two versions, one for lvalues and one for non-lvalues. However, it is important to note that at function resolution time, we don't determine which is called. That is, we always resolve both versions. So this argument can not be used as a parameter.
- `Type* retType` stores the resolved return type of a function. This field is set during function resolution because only at that point can we determine the type.
- `BlockStmt* where` stores the expression that is the where clause for a function. It is evaluated to true or false during function resolution when choosing a candidate. This has language ramifications to our in-order function resolution since we could end up in an infinite recursion if we have to evaluate a where clause by evaluating a function that we have not yet finished evaluating. This results in an error rather than an infinite recursion.
- `BlockStmt* retExprType` stores the specified return type as an expression. The type cannot be resolved from this expression until function resolution time, except for simple cases. This expression is looked at during normalization when normalizing returns (making each function have a single return statement that is the last statement in that function).
- `BlockStmt* body` stores the BlockStmt that is the body of the function. This is where the code is!
- `IntentTag thisTag` stores an intent on the implicit this argument. It is set during parsing before the implicit this argument has been created. Perhaps we could also create the implicit this argument earlier...
- `RetTag retTag` stores the return tag, whether this function is a var function, a param function, a type function, or a regular "value" function.
- `IteratorInfo* iteratorInfo` stores information about the iterator record and class that are created when an iterator function is lowered into a record and a class with methods during the lowerIterators pass. The prototypes are set up during function resolution, and then filled in when iterators are lowered.
- `Symbol* _this` stores a pointer to the implicit this argument for methods. For constructors, `_this` is a VarSymbol that is declared in the body of the constructor.
- `Symbol* _outer` pertains to outer classes. It is set for nested class constructors and replaces `_this`.
- `FnSymbol* instantiatedFrom` is set for functions that have been instantiated from generic functions. This points to the generic function, is only used during function resolution, and can probably be stored in a hash table.
- `SymbolMap substitutions` is a hash table of the substitutions that were made to instantiate this generic function. The keys are arguments in the generic function (original function) and the values are the types or parameters over which these arguments have been instantiated.
- `BlockStmt* instantiationPoint` points to a point in the code that we are using as the instantiation point for function resolution (just like the instantiation point in C++). Constrained generics will not need an instantiation point!
- `Vec<BasicBlock*>* basicBlocks` is built when calling `buildBasicBlocks` on a function. The basic block is stored as a vector of Exprs where these are pointing at the expressions in the function. Basic blocks are used by the dataflow analyses and optimizations. They are not maintained. This probably just does not need to be a field, as it is not updated and is recomputed for each analysis.

- `Vec<CallExpr*>` `calledBy` points to all of the `CallExprs` that call this function. These vectors are computed by `compute_call_sites` (Section 2.6.9), but are not maintained when `CallExprs` are changed. This field could be removed if this function were to build a `Map` from `FnSymbols` to a vector of `CallExprs`.
- `const char*` `userString` is set during parsing so that we can store the prototype of the function and issue error messages to the user that display function candidates just as they are written in the Chapel code.
- `FnSymbol*` `valueFunction` points to a value version of a var function. This is only referenced in function resolution and the pass that culls references. It may be better to use a hash table defined and used by the resolution passes.

2.2.5 EnumSymbol

The `EnumSymbol` class defines an enumeration constant (not an enumeration type). This class has no fields, inheriting all of the information it needs from the `Symbol` class.

2.2.6 LabelSymbol

The `LabelSymbol` class defines a label that can be the target of a `break` or `continue` statement, or a `goto` statement. There are no `goto` statements in Chapel, but there are `goto` statements in our intermediate representation, and they are introduced especially (perhaps only) when iterators are lowered into classes and methods (including an `advance` method). In subsequent calls to the `advance` method, a jump table is implemented with `goto` statements.

The `DefExpr` of the `LabelSymbol` must be declared in a function and is eventually the target of a `goto`. During scope resolution, all `break` and `continue` statements are lowered to normal `goto` statements. The `SymExpr` of a `LabelSymbol` becomes a use of the label. Such uses only occur in `GotoStmts`.

This class has no fields, inheriting all of the information it needs from the `Symbol` class.

2.2.7 ModuleSymbol

Chapel programs are composed of modules and these modules are stored in the `ModuleSymbol`. The fields of `ModuleSymbol` include the following:

- `ModTag modTag` categorizes the type of module as a user module, an internal module, a main module (passed to the compiler explicitly, and thus a candidate for `main`, I think), or a standard module.
- `BlockStmt*` `block` contains the code defined in a module symbol. This is mostly a list of `DefExprs` of functions, types, and global variables.
- `FnSymbol*` `initFn` contains a pointer to the initialize function for a module. This function is called when a module is used at program startup. During the parse pass, it is built and the executable code at module level is moved into this function.
- `const char*` `filename` stores the filename that this module was declared in. Historically, we used to store `filename` alongside `lineno`, but this seemed like overkill given that all of the AST in a module has the same filename. After inlining, this may not be the case, but the current way may be sufficient.
- `Vec<ModuleSymbol*>` `modUseList` is a vector of modules used to determine the order of module initialization. (bradcray r15954)
- `Vec<ModuleSymbol*>` `modUseSet` is a vector of modules used to determine the order of module initialization. Presumably, this field is a set and the previous one is a vector. It is unclear whether these have to be fields. (bradcray r15954)

2.3 Type

Type is basically a class that is associated with TypeSymbols in a one-to-one mapping and stores auxiliary information relevant to the particular kind of type.

The fields in Type are:

- `Vec<Type*> dispatchParents` stores the vector of types that are parents, *e.g.*, a class's superclasses.
- `Vec<Type*> dispatchChildren` stores the vector of types that are children, *e.g.*, a class's subclasses. The inheritance relationship is set up before resolving symbols via scope resolution since we need to know this relationship within methods. Note that the compiler, to date, only supports a single parent though we use a vector here for future support of multiple inheritance.
- `Type* scalarPromotionType` marks the element type for which this type would scalar promote a function.
- `TypeSymbol* symbol` points to the TypeSymbol that represents this type and points to this class via the `type` field.
- `Symbol* defaultValue` points to a symbol that can be used to initialize a variable of this type. If `defaultValue` is NULL (as in the case of a record), then the default constructor is invoked to initialize a variable of this type.
- `FnSymbol* defaultConstructor` points to the default constructor for this type. The default constructor is fairly elaborate in Chapel, having an argument for every field. Even if a type is defined with user-defined constructors, we still create and call the default constructor when resolving the type. However, the default constructor is deprioritized.
- `FnSymbol* defaultTypeConstructor` points to the default type constructor. This name is less than ideal since there are never any type constructors other than the default type constructor. The type constructor is resolved to when specifying a type rather than instantiating an instance of the type. For classes that have fields where the type of the field is omitted, the type constructor will expect a type for this field whereas the default constructor will expect a value. This implements the semantics which allow for a class like `class C { var x; }` to be constructed with a call to the type constructor as `C(int)`.
- `FnSymbol* destructor` points to a destructor. This field is not set until during function resolution and is then used during the pass that inserts calls to the destructors.
- `Vec<FnSymbol*> methods` is a vector of the methods on a type (including primitive types).
- `bool hasGenericDefaults` marks a type if all of the generic fields on the type have default values, since such a type has to be handled with care if the generic version is sought (via a `?`).
- `Type* instantiatedFrom` points to the generic type from which this type was instantiated.
- `SymbolMap substitutions` is a hash table of the substitutions that were made to instantiate this generic type. It is similar to the field on the function symbol.
- `ClassType* refType` points to the reference type if this type is a value type. Reference types are introduced during function resolution because there are no reference types in Chapel, just in the intermediate representation.

The subclasses of Type include the following:

- `PrimitiveType` (Section 2.3.1)
- `EnumType` (Section 2.3.2)
- `ClassType` (Section 2.3.3)

2.3.1 PrimitiveType

The `PrimitiveType` class is used to represent nodes that are primitive types such as integers, reals, strings, C files, a timer type, some auxiliary types for sync and single variables, etc.

This class has no fields, inheriting all of the information it needs from the `Type` class.

2.3.2 EnumType

The EnumeratedType class contains auxiliary information for enumerated types, basically a list of EnumSymbol constants. The fields in EnumType are:

- `AList constants` is a list of DefExprs that define the EnumSymbols for this enumerated type.

2.3.3 ClassType

The ClassType class contains auxiliary information for classes, records, and unions. The fields of ClassType are:

- `ClassTag classTag` is an enumeration to identify whether this ClassType is a class, record, or union. The creative constant name `CLASS_CLASS` suggests this class could have a better name.
- `AList fields` is a list of DefExprs that define the fields in this class, represented as VarSymbols.
- `AList inherits` is a list of Exprs that represent the types from which this type inherits. In the current implementation this is limited to a SymExpr. Inheriting from generics is complicated because we do not instantiate generic types until function resolution, but we need to build up the inheritance hierarchy in time for scope resolution. This poses a challenge, but we should be able to handle it by computing a generic inheritance graph.
- `Symbol* outer` points to the outer class if this is an inner class.

2.4 Expr

The Expr subclasses represents the structure of the program in the AST. There may be multiple pointers to the same instances of Symbol or Type subclasses, but there is only a single instance of a pointer to any Expr subclass, with the notable exception of back pointers such as `parentExpr` and `defPoint`. The back pointer `parentSymbol` is to a symbol and so is not quite as exceptional.

To elaborate, all occurrences of symbols are wrapped by a unique SymExpr instance. So given a program that contains the calls

```
f1(a, b);  
f2(b, c);
```

the AST for this portion of the program will be represented (throughout the entire compilation) by two different instances of CallExpr. Each instance will have the `baseExpr` field point to a distinct instance of the SymExpr class. The `var` fields will point to the FnSymbol instances that represent `f1` and `f2`. Each instance of CallExpr will have the `argList` field point to a list of expressions. Each list will contain two distinct SymExpr instances, for a total of four distinct SymExpr instances. The second SymExpr instance in the first CallExpr will set `var` to point to the Symbol that represents `b`, as will the first SymExpr instance in the second CallExpr. Thus the symbol pointer appears more than once via “forward” pointers, but pointers to Expr subclasses can appear only once via “forward” pointers.

The fields in Expr are:

- `Expr* prev` is a back pointer to the previous expression in a list, if this expression is in a list.
- `Expr* next` is a pointer to the next expression in a list, if this expression is in a list. It may be useful to think of next as a back pointer if we think of the “forward” pointer to the AList as pointing to all of the expressions in the list.
- `AList* list` is a back pointer to the list of expressions to which this expression is a part.
- `Expr* parentExpr` is a back pointer to the parent expression. This is set any time an Expr is pointed to by another Expr. For example, a parent CallExpr node points to its children via `baseExpr` and `argList`. The Expr nodes pointed to via these pointers will set `parentExpr` to point to this parent CallExpr node. The `parentExpr` node can be NULL even if this node is part of the AST. This happens if this Expr node is pointed at directly from a Symbol or Type node. For example, via the `body` field of a FnSymbol node.

- `Symbol* parentSymbol` is a back pointer to the parent symbol. If this is not set, the node is not considered part of the AST. When using `remove` to remove an expression from the AST, all the `Expr` nodes that make up that expression are traversed and their `parentSymbol` back pointers are set to `NULL`. Testing for the existence of `parentSymbol` is sufficient to tell if an `Expr` node is part of the AST. Between passes, nodes that are not part of the AST are garbage collected.

Note that by and large, `prev`, `next`, `list` do not have to be manipulated because inserting and removing `Expr` nodes from the AST should be done via the functions `insertAtHead`, `insertAtTail`, `insertBefore`, `insertAfter`, `replace`, and `remove`. These functions will also update the back pointers `parentExpr` and `parentSymbol`, which should never be set elsewhere. These functions are described in Section 2.6.1.

The subclasses of `Expr` include the following:

- `DefExpr` (Section 2.4.1)
- `SymExpr` (Section 2.4.2)
- `UnresolvedSymExpr` (Section 2.4.3)
- `CallExpr` (Section 2.4.4)
- `NamedExpr` (Section 2.4.5)
- `BlockStmt` (Section 2.4.6)
- `CondStmt` (Section 2.4.7)
- `GotoStmt` (Section 2.4.8)

Note that some of these end with “`Stmt`” while others end with “`Expr`.” This distinction is important as, at least after normalization, we will never get to a `BlockStmt` when traversing the AST from `SymExpr`, `UnresolvedSymExpr`, `CallExpr`, or `NamedExpr` if we do not traverse into `Symbol` nodes unless we are coming from a `DefExpr` node and we do not traverse into `Type` nodes unless we are coming from a `TypeSymbol` node. This is indeed how AST traversal works. See Section 2.5.

2.4.1 DefExpr

Every symbol, with the exception of `rootModule`, is pointed to by a distinct instance of the `DefExpr` class. The forward pointer to this `DefExpr` locates the declaration of this symbol in the AST.

The fields in `DefExpr` are:

- `Symbol* sym` is a pointer to the symbol that this `DefExpr` declares. Note that using the typical terminology of compilers, this is not a `def` for a `VarSymbol` after normalization. A `def` would be marked by a `SymExpr` (an occurrence of the symbol) that is on the left-hand side of an assignment or primitive move, for example.
- `Expr* init` is used to capture an expression that is used to initialize this symbol. For variables and fields, this is the initialization expression in the Chapel code. This is not used for `ArgSymbols` (where the corresponding expression is pointed to directly by the `ArgSymbol` field `defaultExpr`). After normalization, this field should be `NULL` and the AST should be modified so that the functionality is elsewhere. For variables, this means that there will be expressions inserted after this `DefExpr` to do the initialization. For fields, this means that there will be expressions added to a default constructor to do the initialization. The code to complete this transformation is in the `normalize` pass.
- `Expr* exprType` is used to capture an expression that is used to specify the type. Before function resolution, the type is stored as an expression. It could be a simple `SymExpr` if the type is a simple record `R` or a primitive type like `int`, but this need not be the case if the type is specified as a call to a type function or an instantiated type such as `C(int)`. The `exprType` field undergoes similar transformations to the `init` field at the same time (normalization).

2.4.2 SymExpr

Whenever a symbol is used or defined (not declared) in the AST, it is wrapped by an instance of the `SymExpr` class. This makes for a bulky intermediate representation. An optimized intermediate representation might replace the list of expressions in a normalized `CallExpr` by a vector of `Symbols`.

The fields in `SymExpr` are:

- `Symbol* var` is a pointer to the occurrence of the symbol that this `SymExpr` wraps. Later in compilation, it is often useful to categorize these `SymExprs` as either uses or defs. I think there is a slight problem with respect to this categorization as currently involves references.

2.4.3 UnresolvedSymExpr

The `UnresolvedSymExpr` nodes represent names of symbols that have not yet been resolved by either scope resolution (variables, fields, arguments, etc.) or function resolution (functions). This is one of the most recently introduced nodes as it replaces `UnresolvedSymbol`, a subclass of `Symbol`.

The fields in `UnresolvedSymExpr` are:

- `const char* unresolved` is the name of a symbol for which the symbol has not been resolved. Nodes of this type should not exist after function resolution. They are eliminated during scope resolution and function resolution. Comparing this field does not require a string comparison as it is canonicalized (see note on strings in Section 4.1).

2.4.4 CallExpr

The `CallExpr` node represents function calls as well as primitive calls. Primitives are further explained in Section 2.7.

The fields in `CallExpr` are:

- `Expr* baseExpr` is a pointer to the base expression (the expression that evaluates eventually to a `FnSymbol`). After function resolution, this field, if set, will always point to a `SymExpr` that points to a `FnSymbol`. Before function resolution, there are additional possibilities. For example, given the Chapel code `f()()`, where the call `f()` returns an instance of a class with a `this` metho, then the `CallExpr` of `f()` will be pointed to by the `baseExpr` field.
- `AList argList` is a pointer to the expression list that makes up the actual arguments to this `CallExpr`.
- `PrimitiveOp* primitive` is a pointer to the primitive. If this field is set, then `baseExpr` will be `NULL`. It's a one-or-the-other situation.
- `bool partialTag` handles the case involved with resolving the Chapel code given by `x.y()` by marking this case. We have to deal with two situations. The first is where we are invoking a method `y`. The second is where we are invoking a parentheses-less method (possibly a getter for a field) that returns a class instance with a `this` method, and then we are invoking the `this` method. So in the latter case, we invoke two functions, while in the former case, we invoke just one. This code is represented as a nested `CallExpr`, and the inner `CallExpr` is marked with this tag after normalization.
- `bool methodTag` marks basic method invocations, even if they are parentheses-less. For example, with `x.y`, the method call of `y` after normalization will be mark this flag true, but `partialTag` will not be marked true. There are no nested call expressions. It seems that if `partialTag` is set, then `methodTag` is set.
- `bool square` marks function calls that are invoked using square brackets. This is the only distinguishing mark for this case.

2.4.5 NamedExpr

The `NamedExpr` class is used to represent passing an actual to a function via a named argument.

The fields in `NamedExpr` are:

- `const char* name` is the name of the formal argument. Comparing this field does not require a string comparison as it is canonicalized (see note on strings in Section 4.1).
- `Expr* actual` is the actual expression passed to the call.

2.4.6 BlockStmt

The `BlockStmt` class represents a block of code in Chapel, including all loops.

The fields in `BlockStmt` are:

- `BlockTag blockTag` marks the `BlockStmt` with a tag. This field used to be used to distinguish the different types of `BlockStmts` more than it is used now. Most of this functionality has been moved into the `blockInfo` field via a `CallExpr` with a primitive. The remaining tags can be moved there as well, eliminating this field. These tags mark this block in two ways. First, scopeless blocks are marked to indicate that this block should not be considered relevant when completing scope resolution. Second, type blocks are marked to indicate that the code in this block statement should only be used to determine the types of variables (during function resolution). Such blocks are removed after function resolution and the code represented by this block is not executed.
- `AList body` is a list of expressions that form the body of this block, the statements within it.
- `CallExpr* blockInfo` is a call to a primitive that marks this block as special in some way, *e.g.*, a while loop.
- `CallExpr* modUses` is a call of `ModuleSymbols` where each of these `ModuleSymbols` represents a use of a particular module. This list is largely constructed during scope resolution.
- `LabelSymbol* breakLabel` is a pointer to a compiler-inserted `LabelSymbol` that should be the target of a break's `GotoStmt`.
- `LabelSymbol* continueLabel` is a pointer to a compiler-inserted `LabelSymbol` that should be the target of a continue's `GotoStmt`.
- `const char* userLabel` is a label introduced by the Chapel programmer before a loop that can be the target of either a break or a continue's `GotoStmt`.

2.4.7 CondStmt

The `CondStmt` class represents conditional statements (conditional expressions are represented by nested functions that contain conditional statements).

The fields in `CondStmt` are:

- `Expr* condExpr` points to the expression that is evaluated (at compilation time or runtime) to true or false. After normalization, this is a `SymExpr`. If this expression evaluates to a parameter value during function resolution, the conditional statement is folded and removed during function resolution.
- `BlockStmt* thenStmt` points to a block of code that forms the true part of this statement.
- `BlockStmt* elseStmt` points to a block of code that forms the false part of this statement. This block can be `NULL` if there is no else part.

2.4.8 GotoStmt

The `GotoStmt` implements the well-understood, but “harmful” goto statement. Although there are no goto statements in Chapel, there are in our intermediate form. All break and continue statements are transformed into goto statements during scope resolution. Additional goto statements are introduced when lower iterators, as described in Section 3.14.

The fields in `GotoStmt` are:

- `GotoTag gotoTag` marks the type of the goto statement. This distinguishes breaks and continues from normal gotos. After scope resolution, there are no more break or continues. Note that for certain simple cases, we could use break and continue statements throughout compilation, but for more complicated cases, this is impossible since some of the code that must be done for each iteration has to be moved inside the loop and some code that should be done only after the loop completes naturally has to be placed after the loop. Eliminating breaks and continues also, of course, simplifies the number of cases we have to handle later on during compilation. This digression illustrates a philosophical change/difference between the early implementation of the compiler and the later implementation of the compiler. The semantics of the language can be better mapped to the implementation if we use a simple intermediate representation.
- `Expr* label` is often initially an `UnresolvedSymExpr` to a name that should eventually resolve to a `LabelSymbol`. After scope resolution, this field should always point to a `SymExpr` that wraps a `LabelSymbol`.

2.5 Traversing the AST

The AST is a graph that the compiler traverses in a number of ways.

2.5.1 The Core Traversal (With a Figure of the AST Graph)

The core traversal mechanism is implemented by the macro `AST_CHILDREN_CALL`. This macro takes as arguments a `BaseAST` pointer, a C function call, and a list of arguments that should be passed to the C function call. This macro should be used in a recursive function that takes an instance of the `BaseAST` class and any number of arguments. The arguments to the macro must be the `BaseAST` instance, the name of the recursive function, and arguments that should be passed to recursive calls of the function on the child nodes of the AST. A good example of using this macro is in the function `collectDefExprs` which collects all the `DefExprs` that should be traversed from any AST node, for example, a `Fnsymbol`. In this example, we would capture all of the local variable definitions in the function.

The definition of `AST_CHILDREN_CALL` must be updated whenever the AST is modified, either by adding fields or new classes. The implementation of this macro captures the structure of the graph. Notice that when traversing a `CallExpr`, we traverse its `baseExpr` (`baseExpr`) and its list of actuals (`argList`).

The graph structure (omitting leaf nodes like `UnresolvedSymExpr`) is illustrated in Figure 1. Every subclass of `Symbol`, `Type`, and `Expr` is represented if that subclass may not be a leaf in the traversal of the graph. For example, both `SymExpr` and `UnresolvedSymExpr` are leaves when traversing. The `var` field of `SymExpr` is not traversed. (Since we only want to traverse every `Symbol` once, we do it via the `DefExpr`. More than one `SymExpr` node can point to the same `Symbol`, and that is to be expected.)

Although there are cycles in Figure 1 in terms of classes, there are no cycles in terms of instances of classes. That is, even though we can traverse from `DefExpr` to `Fnsymbol` and back to `DefExpr`, the next `DefExpr` will definitely be a different instance, and that in turn will go to a different instance of `Symbol`. Another important point to note is that we only traverse back to either `Symbol` or `Type` from `Expr` nodes via the `sym` field of the `DefExpr` node.

In these data structures, the only other field that points to an `Expr` node is the `parentExpr` field. This is a back pointer and so is not traversed when visiting the nodes via the traversal mechanism.

There are a number of fields that store pointers to symbols. Some of these are back pointers, like `parentSymbol` in `Expr`, but some are not, like `var` in `SymExpr`. Pointers to the same `Symbol` and `Type` instances can appear more than once in this IR, but only once via a `DefExpr`, which roots it in the Chapel program.

A number of collection routines are implemented that use the traversal mechanism to collect a vector of nodes. These routines are sometimes more convenient or more efficient to use, depending on the particular analysis or transform being implemented. The following list captures the current implementation:

```
void collect_ast(BaseAST* ast, Vec<BaseAST*>& asts);
void collect_ast_postorder(BaseAST*, Vec<BaseAST*>& asts);
void collect_top_ast(BaseAST* ast, Vec<BaseAST*>& asts);
void collect_stmts(BaseAST* ast, Vec<Expr*>& stmts);
void collectDefExprs(BaseAST* ast, Vec<DefExpr*>& defExprs);
void collectCallExprs(BaseAST* ast, Vec<CallExpr*>& callExprs);
void collectGotoStmts(BaseAST* ast, Vec<GotoStmt*>& gotoStmts);
void collectSymExprs(BaseAST* ast, Vec<SymExpr*>& symExprs);
void collectSymbols(BaseAST* ast, Vec<Symbol*>& symbols);
void collectFnCalls(BaseAST* ast, Vec<CallExpr*>& calls);
```

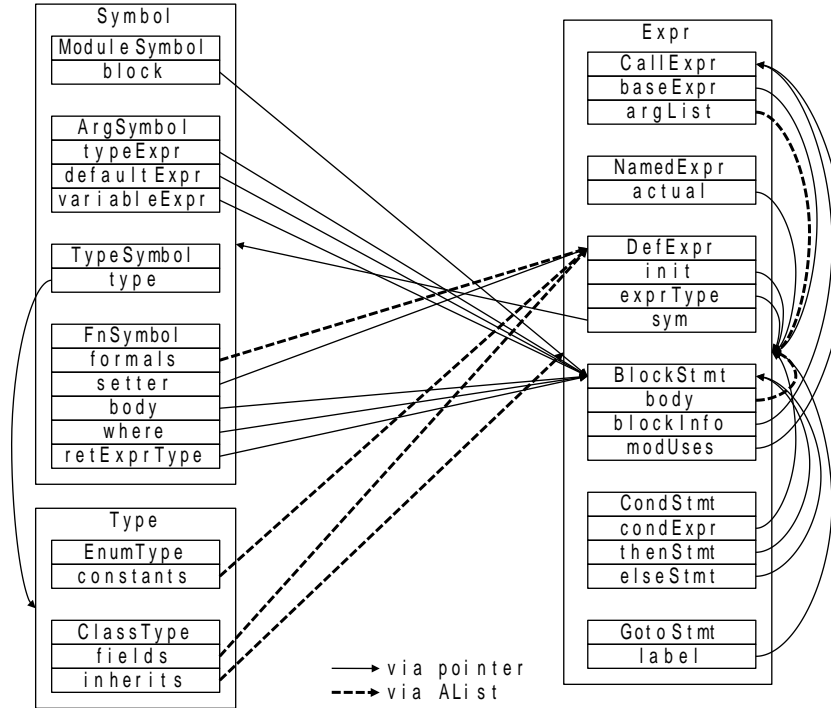


Figure 1. The AST graph structure, as it is traversed, omitting leaf nodes.

Except for `collectFnCalls` and `collect_top_ast`s, these routines traverse the AST in total from a given AST node `ast`, and collect the nodes of a given type into a vector, the second argument. The function `collectFnCalls` collects all `CallExpr`s that are not to primitives. The function `collect_top_ast`s collects all nodes that are “top-level” where that is defined as not traversing into a `Symbol` via a `DefExpr`. Thus before the pass that flattens nested functions, this function is useful if we don’t want to collect the nodes in a nested function when traversing the outer function.

The collect routines are sometimes useful and efficient, but for time-critical sections of code, it is sometimes better to define a recursive function that calls the traversal macro `AST_CHILDREN_CALL` directly.

2.5.2 The `getFirstExpr` and `getNextExpr` Methods

The stylized loop `for_exprs_postorder` uses the functions

```
Expr* getFirstExpr(Expr* expr);
Expr* getNextExpr(Expr* expr);
```

to iterate over an expression. This iteration technique is only used in function resolution to date.

The function `getFirstExpr` returns the first expression that would be returned in a postorder traversal of the `Expr` node that is passed to this function. Then given any `Expr` node in this traversal, the function `getNextExpr` returns the next expression in the postorder traversal. This functionality can only be used on “top-level” `Expr` nodes, *i.e.*, `Expr` nodes that have a null `parentExpr` field. Otherwise, we would traverse up into the parent.

This style of iteration is used in function resolution when walking through a `BlockStmt` and resolving types and function calls.

2.6 AST Functions and Methods

This section details implemented functions and methods pertaining to the AST.

2.6.1 Methods for Insertion, Replacement, and Removal

For changing the IR, these methods should always be used as they update back pointers and enforce some constraints.

```
void FnSymbol::insertAtHead(Expr* ast);
void FnSymbol::insertAtTail(Expr* ast);
void CallExpr::insertAtHead(BaseAST* ast);
void CallExpr::insertAtTail(BaseAST* ast);
void BlockStmt::insertAtHead(Expr* ast);
void BlockStmt::insertAtTail(Expr* ast);
```

The `insertAtHead` and `insertAtTail` function insert new `Expr` nodes at the beginning or end of a list of `Exprs`. When using these methods on `FnSymbol`, the list is the `body` field in the `BlockStmt::body` field of the `FnSymbol`. When using these methods on `BlockStmt`, the list is the `body` field of the `BlockStmt`.

When using these methods on `CallExpr`, the list is the `argList` field, the list of actual expressions passed to the `CallExpr`. Notice that these methods takes any `BaseAST`, rather than `Expr`. This allows `Symbol` nodes to be passed to these methods directly. The methods insert a `SymExpr` around such `Symbol` nodes. This is important for keeping the code more readable. For example, one can write

```
call->insertAtHead(gMethodToken)
```

instead of

```
call->insertAtHead(new SymExpr(gMethodToken))
```

Such simplifications are important for readability. They also help move towards the day when the `SymExpr` node can be removed.

```
void Expr::insertBefore(Expr* new_ast);
void Expr::insertAfter(Expr* new_ast);
```

These methods insert new `Expr` nodes before or after another `Expr` node that is already in a list. This could be an actual in a `CallExpr`, a statement in a `BlockStmt`, a formal argument in a function, a field in a class, ...

```
void Expr::replace(Expr* new_ast);
```

This method replaces one expression that is in the AST (has a `parentSymbol`) with another one that is not in the AST. It is implemented on each node type via the internal helper method called `replaceChild`. You cannot call `replace` a `NULL` field. To handle that case, we sometimes use the `insert_help` function discussed below.

```
Expr* Expr::remove(void);
```

This method removes an `Expr` node from the AST and updates all forward and backward pointers in the removed expression and what is left. `Expr` nodes can be removed from lists. They can also be removed even if they are not in lists. In this case, the field that was pointing to this expression is set to `NULL`. For example, calling `block->blockInfo->remove()` will remove the `CallExpr` node pointed to by the `blockInfo` field from the AST (making all of its `parentSymbol` fields `NULL`, etc.). In addition, the `blockInfo` field will be set to `NULL`. This uses the same `replaceChild` method to get to this parent pointer from within the call to this method.

```
void FnSymbol::insertBeforeReturn(Expr* ast);
void FnSymbol::insertBeforeReturnAfterLabel(Expr* ast);
```

These special-purpose insert routines also update back pointers, etc. The methods `insertBeforeReturn` and `insertBeforeReturnAfterLabel` insert code in a function just before the return statement. The “after label” modifier is important if want this code to execute even if we jump to this point in the code to exit (as is done after the return statements are normalized by the `normalize` pass).


```

void FnSymbol::insertFormalAtHead(BaseAST* ast);
void FnSymbol::insertFormalAtTail(BaseAST* ast);

```

These methods insert formal arguments at the beginning or end of the formal arguments list of a function. The formal arguments list is a list of DefExpr nodes. To improve readability, an ArgSymbol may be passed to these methods. The DefExpr node is then inserted within these methods.

```

void insert_help(BaseAST* ast, Expr* parentExpr, Symbol* parentSymbol);
void remove_help(BaseAST* ast, int dummy=0); // dummy is never used
void parent_insert_help(BaseAST* parent, Expr* ast);
void sibling_insert_help(BaseAST* sibling, BaseAST* ast);

```

These functions are used to set the back pointers after an insertion, removal, or replacement. When replacing NULL fields, it is sometimes useful to just assign them the nodes that are not yet in the AST, and then call the insert_help function directly. In general though, these are not meant to be called outside of the methods discussed above.

2.6.2 Constructors and new_Expr

As with the insertion methods discussed, constructors will also automatically insert SymExpr wrapper nodes to make for more readable code. The following is a list of such simplifications:

- The DefExpr constructor will automatically insert SymExpr nodes around symbols passed to initialize the `init` and `exprType` fields.
- The CallExpr constructor will automatically insert a SymExpr node around a Symbol passed to initialize the `baseExpr` field.
- The CallExpr constructor can take up to four arguments to fill out the `argList` field. If these arguments are Symbol nodes, they will be wrapped by SymExpr nodes.
- The CondStmt constructor can take non-block expressions to initialize the `thenStmt` and `elseStmt` blocks, in which case a new BlockStmt node will be created to wrap these expressions. In addition, if the BlockStmt nodes passed to this constructor are not regular (have `blockInfo` or `blockTag` set, then a new BlockStmt will be created.

The CallExpr constructor deserves extra mention. It is overloaded so that the first argument can either be a Symbol or an Expr (to initialize the `baseExpr` field, or it can be a PrimitiveOp or PrimitiveTag (to initialize the `primitive` field, or it can be a character strings (to initialize the `baseExpr` field as an unresolved call using an UnresolvedSymExpr).

A relatively new way of creating Expr nodes is to use the `new_Expr` functions.

```

Expr* new_Expr(const char* format, ...);
Expr* new_Expr(const char* format, va_list vl);

```

These functions take a format string and a variable list of arguments and builds up an Expr node. It is meant to be used early in compilation.

Primitives can be created by delimiting the name of a primitive in the format string followed by parentheses, with the arguments to the primitive in the parentheses. Unresolved function calls (to be resolved in function resolution) can be specified by writing the name of the function followed by parentheses, with the arguments to the function in the parentheses. BlockStmts can be created by delimiting the format string with curly braces. BLOCK_TYPE BlockStmts can be created by delimiting the format string with curly braces where the word TYPE appears immediately inside the curly brace.

Expressions and Symbols can be added to the variable list of arguments by using the specifiers %E and %S respectively.

For example, the code

```

new_Expr("'move' (%S, foo(%S))", x, gFalse))

```

is equivalent to


```
new CallExpr(PRIM_MOVE, T2, new CallExpr("foo", gFalse))
```

This is often more readable for long sequences of expressions or statements, but not much rewriting has been done to date.

The extensions to the `insertAtHead` and `insertAtTail` methods, as given by

```
void FnSymbol::insertAtHead(const char* format, ...);
void FnSymbol::insertAtTail(const char* format, ...);
void BlockStmt::insertAtHead(const char* format, ...);
void BlockStmt::insertAtTail(const char* format, ...);
```

make use of the `new_Expr` functionality above to avoid the explicit call to `new_Expr`.

2.6.3 BaseAST Functions and Methods

```
virtual void BaseAST::verify() = 0;
```

This method is called between passes to verify that this AST node, which is still in the AST, is valid. It checks to make sure assumptions are valid.

```
virtual BaseAST* BaseAST::copy(SymbolMap* map = NULL, bool internal = false) = 0;
```

This method is used to copy an `Expr`, `Symbol`, or `Type` node. It is implemented in such a way so that if you copy a `BlockStmt` that contains a `DefExpr` of some `Symbol` node, then that `Symbol` node will be copied and all uses of it (via `SymExpr`, etc.) will be updated to point to the new one. This update is done via a call to `update_symbols` as described below. The map that is used by `update_symbols` can be passed to the copy function. If none is passed, a new one will be created. This map, from original symbols to copied symbols, will be constructed during the copy process. It is sometimes useful to capture this map when calling `copy`.

The argument `internal` should never be passed, except internal to the implementation of `copy`. The implementation of `copy` on a particular node is implemented via the method `copyInner` which is also called by all recursive copies via the macro `COPY_INT`.

```
virtual void BaseAST::codegen(FILE* outfile) = 0;
```

This method is used to implement the code generation pass. Although originally designed for before the AST is normalized (and there was more nesting of expressions as well as more `Expr` nodes), it still works fairly well today. This code might benefit from a revamp with the assumption of normalization, unless there remains a desire to unnormalize so as to generate C expressions.

```
ModuleSymbol* BaseAST::getModule();
FnSymbol* BaseAST::getFunction();
```

These methods return the module/function that a given AST node resides in, by tracking back up the back pointers. In many cases within the compiler, we make the assumption that `parentSymbol` points to a function and avoid the need to call `getFunction`.

```
virtual Type* BaseAST::typeInfo(void) = 0;
Type* BaseAST::getValType();
Type* BaseAST::getRefType();
Type* BaseAST::getWideRefType();
```

The method `typeInfo` returns the type of any expression or symbol, though before function resolution, this is likely to be the unresolved type `dtUnknown`. The method `getValType` returns the value type of the expression or symbol. That is, if the type evaluates to a reference or wide reference type, the value type is returned. If the type evaluates to a value type, it is returned. The methods `getRefType` and `getWideRefType` return reference or wide reference types.

```
void update_symbols(BaseAST* ast, SymbolMap* map);
```

This function replaces all occurrences of the key symbols in `map` with the value symbols in `map`.

2.6.4 Symbol Functions and Methods

```
Symbol* FnSymbol::getReturnSymbol();
```

This method can be used after normalization to find the symbol that a function returns. After normalization, there is only one such symbol, and it is returned by the last statement in the function.

```
int FnSymbol::numFormals();
ArgSymbol* FnSymbol::getFormal(int i);
```

These methods return the number of formal arguments to a function and the *i*th formal argument.

```
VarSymbol *new_StringSymbol(const char *s);
VarSymbol *new_BoolSymbol(bool b, IF1_bool_type size=BOOL_SIZE_SYS);
VarSymbol *new_IntSymbol(int64_t b, IF1_int_type size=INT_SIZE_32);
VarSymbol *new_UIntSymbol(uint64_t b, IF1_int_type size=INT_SIZE_32);
VarSymbol *new_RealSymbol(const char *n, long double b, IF1_float_type size=FLOAT_SIZE_64);
VarSymbol *new_ImagSymbol(const char *n, long double b, IF1_float_type size=FLOAT_SIZE_64);
VarSymbol *new_ComplexSymbol(const char *n, long double r, long double i,
                             IF1_complex_type size=COMPLEX_SIZE_128);
VarSymbol *new_ImmediateSymbol(Immediate *imm);
```

These functions build new symbols to represent immediate or literal values. A cache is used so that we never have two different symbols represent identical literal values.

2.6.5 Type Functions and Methods

```
Symbol* ClassType::getField(const char* name, bool fatal=true);
Symbol* ClassType::getField(int i);
```

These methods return a field in a class, record, or union either by matching a string name or finding the *i*th field by declaration order. The argument *fatal* can be set to false to return NULL rather than issue an internal error.

```
bool is_bool_type(Type*);
bool is_int_type(Type*);
bool is_uint_type(Type*);
bool is_real_type(Type*);
bool is_imag_type(Type*);
bool is_complex_type(Type*);
bool is_enum_type(Type*);
bool isClass(Type* t);
bool isRecord(Type* t);
bool isUnion(Type* t);
bool isReferenceType(Type* t);
```

These functions return true if the *Type* node represents a Chapel type of the named category. In addition, a macro *is_arithmetic_type* wraps calls to the numeric type query functions above.

```
int get_width(Type*);
```

This function returns the number of bits in a type.

2.6.6 Expr Functions and Methods

```
Expr* getStmtExpr();
```

This method returns a statement-level expression given an expression nested in another. A statement-level expression is an expression that is a statement or an expression whose parent is a statement. A statement is a BlockStmt, a GotoStmt, or a CondStmt. Back in the day, the Chapel IR used to distinguish between expressions and statements, even going so far as having a statement called ExprStmt. When working on the compiler, it is worth keeping in mind that our intermediate representation has been greatly simplified over time. Thus some code that may seem like it should have been simplified from when it was first written could not have been, but it is probably worthwhile simplifying now. That is to say, if something seems more complicated than it needs to be, it is quite possible that it is more complicated than it needs to be now. Therefore, simplify!

```
FnSymbol* CallExpr::isResolved(void);
```

This method returns the FnSymbol that a CallExpr has been resolved to. This method tends to be called extensively after function resolution. For primitives and unresolved symbols, NULL is returned.

```
bool CallExpr::isNamed(const char*);
```

This method returns true if the name of the function call (be it resolved or unresolved) matches the string argument.

```
int CallExpr::numActuals();  
Expr* CallExpr::get(int index);
```

The method numActuals returns the number of actual expressions passed to this call. The method get returns the indexth actual expression passed to this call.

```
bool CallExpr::isPrimitive(PrimitiveTag primitiveTag);  
bool CallExpr::isPrimitive(const char* primitiveName);
```

These methods return true if this call is a primitive that matches the enumeration or string argument.

```
bool get_int(Expr *e, long *i);  
bool get_uint(Expr *e, unsigned long *i);  
bool get_string(Expr *e, const char **s);  
const char* get_string(Expr* e);  
VarSymbol *get_constant(Expr *e);
```

These functions evaluate an expression and return (or return in the reference argument) the value of the compile-time constant that the expression evaluates to. In the event that the expression is not a compile-time constant, the function either returns false (get_int, get_uint, the first get_string), issues an error (the second get_string), or returns NULL (get_constant).

2.6.7 Stmt (Expr) Functions and Methods

```
bool BlockStmt::isLoop(void);
```

This method returns true if this BlockStmt node is a loop. This method is safe during the scopeResolve pass and before. Use of this method should probably be discouraged at or after function resolution when there are more loop types and basic blocks can be used for such control flow analyses.

```
int BlockStmt::length(void);
```

This method returns the number of Expr nodes top-level to a block. How long is this block? This may not be very applicable as all expressions are counted (including DefExpr nodes) so what exactly is this doing? This method is used during parsing and also appears to be used during the fast and short on-statement optimization.

```
Expr* CondStmt::fold_cond_stmt();
```

This method folds a conditional statement if the expression that evaluates to true or false can do so at compilation time. This is used during function resolution to fold parameter conditionals. It is also used after function resolution if new parameter conditionals are introduced. Such new parameter conditionals show up if a value is changed to true or false after function resolution. Currently, this definitely does happen on occasion.

2.6.8 Miscellaneous AST Utility Functions

```
ArgSymbol* actual_to_formal( Expr *a);
Expr* formal_to_actual(CallExpr* call, Symbol* formal);
```

The function `actual_to_formal` finds the formal argument from an actual argument (expression). The function `formal_to_actual` finds the actual argument in the specified call from a formal argument.

```
void subSymbol(BaseAST* ast, Symbol* oldSym, Symbol* newSym);
```

This function replaces all occurrences of symbol `oldSym` in `ast` with `newSym` by traversing `ast`. This is a special case of `update_symbols` which takes a map of old symbols to new symbols.

```
BlockStmt* getVisibilityBlock(Expr* expr);
```

This function returns the innermost block that could contain definitions that may be resolved to during scope resolution or function resolution. From this block, it is a matter of searching into outer blocks, or module blocks used by this block or outer blocks, etc.

```
void reset_line_info(BaseAST* baseAST, int lineno);
```

This helper function resets the line number in `baseAST` and all AST nodes traversed from this node.

2.6.9 Compilation Utility Routines

```
void compute_call_sites();
```

This function builds the call graph for the entire program represented by the AST. Each `FnSymbol` has a field `calledBy` that is a vector of `CallExprs`. After this function is called, these vectors are filled with every `CallExpr` that may call this function. This includes dynamically dispatched calls. This function can be used more than once; the vectors are cleared at the beginning of the call.

```
void buildDefUseMaps(Map<Symbol*,Vec<SymExpr*>>& defMap,
                    Map<Symbol*,Vec<SymExpr*>>& useMap);
void buildDefUseMaps(FnSymbol* fn,
                    Map<Symbol*,Vec<SymExpr*>>& defMap,
                    Map<Symbol*,Vec<SymExpr*>>& useMap);
void buildDefUseMaps(Vec<Symbol*>& symSet,
                    Map<Symbol*,Vec<SymExpr*>>& defMap,
                    Map<Symbol*,Vec<SymExpr*>>& useMap);
void buildDefUseMaps(Vec<Symbol*>& symSet,
                    Vec<SymExpr*>& symExprs,
                    Map<Symbol*,Vec<SymExpr*>>& defMap,
                    Map<Symbol*,Vec<SymExpr*>>& useMap);
void freeDefUseMaps(Map<Symbol*,Vec<SymExpr*>>& defMap,
                    Map<Symbol*,Vec<SymExpr*>>& useMap);
void addDef(Map<Symbol*,Vec<SymExpr*>>& defMap, SymExpr* def);
void addUse(Map<Symbol*,Vec<SymExpr*>>& useMap, SymExpr* use);
```

The overloaded functions called `buildDefUseMaps` build def and use maps for all symbols in the entire program, for the variables in a particular function, for the symbols in a set, or for the symbols in a set and the defs and uses in a vector (respectively).

The defs and uses are store in vectors of `SymExprs`. Both defs and uses are `SymExprs`. Whether a particular `SymExpr` is one or the other depends on where it occurs, *e.g.*, the left or right hand side of a move primitive.

The function `freeDefUseMaps` frees these maps.

The functions `addDef` and `addUse` can be used to add defs and uses to these maps incrementally, to avoid recomputing when changing the AST. These maps are not maintained in general.

Two stylized loop macros allow for iteration over the defs and uses. The macro `for_defs(def, defMap, sym)` declares a new `SymExpr` for the `def` argument and iterates over the vector of definitions for the symbol `sym` given by the map `defMap`. The macro `for_uses(use, useMap, sym)` is comparable, but for uses.

```
void collectSymbolSetSymExprVec(BaseAST* ast,
                               Vec<Symbol*>& symSet,
                               Vec<SymExpr*>& symExprs);
```

This function traverses the `ast` and fills a set of `Symbol` nodes and a vector of `SymExpr` nodes with all encountered `Symbol` and `SymExpr` nodes. This set and vector can be passed to one of the above functions that build def and use maps. Sometimes it is useful to compute these separately so that they can be reused (rather than have them computed in a different function that builds def and use maps).

```
void buildDefUseSets(Vec<Symbol*>& syms,
                   FnSymbol* fn,
                   Vec<SymExpr*>& defSet,
                   Vec<SymExpr*>& useSet);
```

This function builds the def and use sets for a vector of symbols as they occur in the specified function `fn`. The set `defSet` contains all the defs. The set `useSet` contains all the uses. This data structure is used during copy propagation, reaching definitions analysis, and live variable analysis to make it fast to determine whether a given `SymExpr` is a def or a use or both.

2.6.10 Optimizations and Analyses

```
void collapseBlocks(BlockStmt* block);
```

This function collapses all blocks that are unnecessary, including blocks that may introduce new scopes.

```
void removeUnnecessaryGotos(FnSymbol* fn);
void removeUnusedLabels(FnSymbol* fn);
```

The function `removeUnnecessaryGotos` removes unnecessary `goto` statements, *i.e.*, `goto` statements that immediately precede the label to which they go to. The function `removeUnusedLabels` removes labels that are not targeted by any `goto` statements.

```
void localCopyPropagation(FnSymbol* fn);
void globalCopyPropagation(FnSymbol* fn);
```

These functions implement local and global copy propagation. This code is based on the algorithm described in “Advanced Compiler Design and Implementation” by Steven Muchnick.

```
void eliminateSingleAssignmentReference(Map<Symbol*,Vec<SymExpr*>>& defMap,
                                         Map<Symbol*,Vec<SymExpr*>>& useMap,
                                         Symbol* var);
void singleAssignmentRefPropagation(FnSymbol* fn);
```

These functions attempt to eliminate references in a similar fashion to the way copy propagation tries to eliminate variables, except this function limits itself to references that are assigned only once.

```
void liveVariableAnalysis(FnSymbol* fn,
    Vec<Symbol*>& locals,
    Map<Symbol*,int>& localID,
    Vec<SymExpr*>& useSet,
    Vec<SymExpr*>& defSet,
    Vec<BitVec*>& OUT);
```

This function computes a live variable analysis on a function. The code is based on the algorithm described in “Advanced Compiler Design and Implementation” by Steven Muchnick. Live variable analysis is used when lowering iterators—The iterator class only needs to store the local variables in the iterator that are live.

```
void buildDefUseChains(FnSymbol* fn,
    Map<SymExpr*,Vec<SymExpr*>>& DU,
    Map<SymExpr*,Vec<SymExpr*>>& UD);
void freeDefUseChains(Map<SymExpr*,Vec<SymExpr*>>& DU,
    Map<SymExpr*,Vec<SymExpr*>>& UD);
void reachingDefinitionsAnalysis(FnSymbol* fn,
    Vec<SymExpr*>& defs,
    Map<SymExpr*,int>& defMap,
    Vec<SymExpr*>& useSet,
    Vec<SymExpr*>& defSet,
    Vec<BitVec*>& IN);
```

The first two functions build and free def-use (DU) and use-def (UD) chains. This code is based on the algorithm described in “Advanced Compiler Design and Implementation” by Steven Muchnick. To build these chains, the compiler completes a reaching definitions analysis, the code of which is based on the algorithm described in this same book.

```
void deadVariableElimination(FnSymbol* fn);
void deadExpressionElimination(FnSymbol* fn);
void deadCodeElimination(FnSymbol* fn);
```

These functions eliminate dead code. The function `deadVariableElimination` eliminates variables that are not used and otherwise unnecessary. The function `deadExpressionElimination` eliminates expressions that do not need to be evaluated because they neither produce values nor have side effects. For example, a statement consisting only of a `SymExpr` never needs to be evaluated—why codegen `x`;

The function `deadCodeElimination` is based on the algorithm described in “Advanced Compiler Design and Implementation” by Steven Muchnick. This uses def-use and use-def chains. This function calls the other two.

2.7 Primitives

Primitives implement the primitive functionality that the Chapel compiler can reason about. For example, the common primitive `PRIM_MOVE` or “move” implements C-level assignment. The left-hand side, or first actual, is the lvalue and must be a `SymExpr`. The right-hand side can be another `SymExpr` node or a `CallExpr` node. After normalization, this is the only case of nested call expressions.

2.8 Pragmas

Pragmas in internal Chapel code are translated into flags in the compiler. Flags are stored on symbols in a bit vector. They are defined by a large enumerated type where each enumeration constant starts with the prefix `FLAG_`. The name of the flag should match the name of the pragma where underscores in the flag are changed to spaces in the pragma and uppercase letters in the flag are changed to lowercase letters in the pragma.

Flags are an important aspect of the Chapel implementation and the compiler often treats constructs specially based on flags. The flags are manipulated via the following methods on `Symbol`:

```

bool Symbol::hasFlag(Flag flag);
void Symbol::addFlag(Flag flag);
void Symbol::addFlags(Vec<const char*>* str);
void Symbol::copyFlags(Symbol* other);
void Symbol::removeFlag(Flag flag);

```

These functions check to see if a flag applies to a symbol (the symbol has a flag), add flags to a symbol (`addFlags` takes the flags as pragmas (strings) and is used in the parser), copy flags (also used in the parser), and remove flags from a symbol.

3 Passes

The compiler is organized as a set of passes, each of which is a function. Compilation proceeds by calling each of these functions in turn. Between each pass, a number of verification and cleanup tasks are completed, as explained in Section 3.34.

As the passes proceed, the assumptions about what the AST looks like change. The passes that impact these assumptions are:

- `parse`
- `cleanup`
- `scopeResolve`
- `flattenClasses`
- `normalize`
- `resolve`
- `flattenFunctions`
- `cullOverReferences`
- `lowerIterators`
- `parallel`

Notice these are all early passes. There are two major shifts and it is sometimes useful to think of the three phases of compilation. These phases are “before normalization” (before `normalize`), “after normalization” (after `normalize`), and “after function resolution” (after `parallel`). The descriptions of the changing assumptions are in the subsections related to the passes where the assumptions change.

3.1 `parse`

The `parse` pass reads the Chapel code from a file and builds up the AST (alternatively called the IR or intermediate representation). This pass issues syntax errors.

There are relatively few assumptions that can be made about the AST at this point in compilation, other than those enforced by the type system.

3.2 `checkParsed`

This pass checks the semantics of the Chapel code as parsed. Although we can’t do all the checks now, we can do a few. This pass, `checkNormalized` and `checkResolved` were meant to complete all of the necessary checks. However, many user errors are issued from other passes, and that was the case before and after these passes were introduced.

This pass checks the following:

- Are explicit argument names repeated in the same function call?
- Do any variables omit both a type and an initializer?

- Are parameters left uninitialized?
- Are configuration variables, constants, and parameters not at module scope?
- Do 'this' and 'these' methods omit parentheses?
- Is a return statement outside of a function?
- Do some returns in a function not return a value while others do?

3.3 cleanup

This pass is meant to clean up the AST as it is parsed. It mostly does things that can't be done while parsing but can be done before scope resolution. It does the following:

- Moves all function definitions that may appear in CallExprs to statement level (as defined by `getStmtExpr`). Prior to this, function definitions may appear in call statements for the nested functions that the compiler inserts (such as the functions for conditional expressions). With support for anonymous functions, this functionality could be generalized. The iterators that are built up for sequential and parallel loop expressions are special-cased so that they are pulled further out.
- Removes all scopeless blocks, blocks that simply group multiple statements together so that the parser can return multiple statements when parsing a single Chapel statement.
- Destructures tuples used on the left-hand side of an assignment statement recursively, transforming the single assignment into one assignment per tuple component using calls to the tuple access function.
- Move primary methods out of the type so that they appear as siblings to the type.
- Change cast expressions in where clauses to be true or false expressions involving the `ISSUBTYPE` primitive.

3.4 scopeResolve

The primary purpose of this pass is to resolve occurrences of names to symbols for variables, arguments, types, etc. After this pass, the only remaining `UnresolvedSymExpr` nodes are for unresolved functions, which are resolved during function resolution.

Scope resolution is a two step process. Between these steps, a number of miscellaneous actions are taken.

In the first step, the compiler constructs a symbol table that roots the declarations of symbols to a particular point in the AST (either a `BlockStmt`, a `TypeSymbol` (for fields and methods), or a `FnSymbol` (for arguments or query identifiers). Multiple definition errors are encountered at this point. The function `lookup` looks for a name in a scope and returns a symbol. Module uses are taken into account, including cyclic module uses.

Then before the second step, the following actions are taken:

- All "use" statements are analyzed. The modules being used have to be looked up before we lookup general names, because this can impact the mapping of names to symbols.
- The class hierarchy is constructed (`dispatchParents` and `dispatchChildren`). This has to be done before we can lookup names of symbols, since we need to be able to find occurrences of fields in super classes from within methods.
- All named arguments using the array alias syntax are marked.
- Constructors and type constructors are built. The default constructor is used during function resolution even if user-defined constructors will shadow it. This helps to determine the types of fields. The type constructor is used to resolve types. Instantiating a generic type looks like a call to the constructor in terms of the generic arguments, and the type constructor is used to resolve such types using the mechanisms of function resolution.
- The type associated with methods is resolved before resolving all other symbols because we need to know the type associated with a method for resolving fields.

- The labels associated with goto statements are resolved. In addition, unlabeled break and continue statements are resolved to the label associated with the innermost loop.

In the second step, the compiler tries to resolve `UnresolvedSymExpr` nodes by looking up their name (`unresolved`) in the symbol table that was created in the first step. Function calls without parentheses are handled now, but function calls with arguments and method invocations are handled during function resolution following a different algorithm that allows for overloading.

Lastly, enumerated types are resolved and the symbol table is destroyed.

3.5 `flattenClasses`

This pass moves the declarations of nested classes to module scope. This functionality could probably be slipped into the `normalize` pass. It just isn't doing much of anything.

3.6 `normalize`

This pass transforms the AST into a normalized form. The transformations can be grouped into two sets, those transforms done in the following function and those not:

```
void normalize(BaseAST* base);
```

This particular function is called on functions that the compiler creates after this pass but before function resolution completes, namely wrapper functions created during function resolution and default functions created during `buildDefaultFunctions`.

The `normalize` function makes the following transformations:

1. Handle syntactic sugar for distributions.
2. Normalize return statements so that each function has a single return statement as the last statement in the function. Other return statements are replaced by a use of the “move” primitive to assign the result to a single return symbol followed by a goto statement to a label immediately before the return statement at the end of the function.
3. Lower `DefExpr` statements for variables so that the `init` and `exprType` fields become `nil`, and “move” primitives are used after the `DefExpr` statement to initialize the value as appropriate.
4. Lower member invocations so that a method call is transformed into a `CallExpr` node with a `this` argument, and a method token argument.
5. Insert temporaries to avoid any nested `CallExpr` nodes. The single exception, and an important and widespread case, is a “move” primitive with a `CallExpr` as the right-hand side argument.
6. Insert a “move” primitive around user-level invocations of assignment.
7. Call constructors for class type instantiations.

In addition to calling the `normalize` function on the AST, the `normalize` pass makes the following transformations:

1. Identify iterator functions based on the existence of yield statements.
2. Replace “delete” primitives with calls to `chpl_destroy`.
3. Lower array formal arguments to possibly generic array arguments with potential reindexing done within the function.
4. Clone methods on the complex type to work on all complex sizes.
5. Lower query identifiers in formal arguments and clone functions on queried primitive types to all sizes.
6. Identify methods of the name type and change them into constructors.
7. Call the `normalize` function on the entire AST.

8. Check for invalid configuration parameters.
9. Check for “use before def” errors.
10. Move functions out of a module’s initialize function since these are global functions.
11. Insert “use” scoping for functions that are called via an explicit module, *e.g.*, `M.foo()`.
12. Check for invalid use of “new” keyword.
13. Insert functions around any statement-level SymExpr to ensure that sync and/or single variables are read.
14. Resolve simple argument types and set argument types for arguments specified with default values but not types. The type is taken from the default value.
15. Check for invalid destructor uses.
16. Check for invalid operator methods.

3.7 checkNormalized

This semantic checks pass flags errors if any of the following statements are true:

- Arguments to iterators have intents.
- Iterators return types or parameters.
- A constructor accesses a member in the type or default value for any argument.

3.8 buildDefaultFunctions

This pass creates many methods, functions, and operators to support functionality, mainly on records and classes, if the user does not supply this functionality, *e.g.*, a `writeThis` method for a class. In many places, the `normalize` function is called on the functions that are created here. We don’t do this before normalization because it is more difficult to determine if the user has created a function that should replace the compiler-generated default.

In addition, this pass creates a *main* function if none exists.

The following functions, methods, and/or operators are created (some of which may not be overwritten by the programmer according to Chapel semantics):

- Getter methods for every field or enumeration constant.
- Destructors for every class and record.
- Default read functions for enumerations, classes, and records.
- Default cast to string functions for enumerations.
- Default write functions for classes and records.
- Equality and inequality operators for records.
- Assignment functions for records, unions, and enumerations.
- Cast functions for records and enumerations.
- Copy functions for records (for use when initializing variables without a specified type).
- Hash functions for records (for use with associative domains).
- Functions that return a tuple of enumeration constants for enumerations.

3.9 resolve

The main purpose of this pass is to resolve function calls and types. After this pass, there are no more `UnresolvedSymExpr` nodes.

This pass is complicated by the vagaries of the Chapel language, for better and for worse.

The algorithm for this pass is summarized as follows:

1. Mark generic functions as generic and identify arguments that are generic even though the generic type of the argument has default values for all of its generic fields, *e.g.*, `range(?)`.
2. Call `resolveFns` (See Section 3.9.1.) on `chpl_main`, the main function in the program, either generated by the compiler during `buildDefaultFunctions` or written by the Chapel programmer.
3. If we are compiling the runtime, which has entries other than through `main`, the resolve the formal arguments of these entry points via `resolveFormals` and the functions via `resolveFns`.
4. Build the virtual method table and resolve methods that are only dynamically dispatched. This work is done in a loop to handle the case where a method that is only resolved via a dynamic dispatch creates a new instantiation of a generic class and this generic class results in a new method that needs to be added to the virtual method table, and then resolving this new method results in another generic class, etc.
5. Resolve calls to `chpl__convertValueToRuntimeType` on any runtime type. Resolve `autoCopy` and `autoDestroy` calls. Resolve record initialization.
6. Replace calls to potentially dynamically dispatched methods with calls via the virtual method table or with conditionals (if there are fewer than some threshold number of such calls).
7. Clean-up the AST by inserting return temporaries for functions that return values, but which are not captured, and by pruning the AST, including eliminating unused functions and types, eliminating method tokens, etc.

3.9.1 resolveFns (fn)

This function resolves the interior of `FnSymbol` node `fn`. When resolving, recursive calls are made to `resolveFns` in order to resolve other functions that are called. Since `resolveFns` is recursive, we use a set to avoid resolving a function that we have already resolved or started to resolve. The algorithmic flow of `resolveFns (fn)` is as follows:

1. If `fn` has already been resolved or is being resolved, stop.
2. If this is a var function, build a value function and resolve the value function with the setter argument set to false. Set the setter argument to true and continue resolving the var function.
3. Insert temporaries for formal arguments to implement copy semantics.
4. Call `resolveBlock` (See Section 3.9.2.) to resolve the body of this function.
5. Determine the return type of this function by examining all of the writes to the returned symbol. If there is no best type (a type the others can dispatch to), flag an error. Otherwise set the `retType` field of `fn`.
6. Insert casts on any `MOVE` primitive that requires it. This is almost certainly done to handle the case of a function returning different types of values in different return statements. Now that the return type has been determined, the appropriate cast can be inserted.
7. If `fn` is resolved as an iterator, prototype iterator records, classes, and methods. (See Section 3.9.9.)
8. If `fn` is resolved as a type constructor, identify the type. Resolve the type constructors of the parent type and all field types. Also resolve the default constructor with no arguments and the destructor.

3.9.2 `resolveBlock(fn)`

This function resolves the body of `FnSymbol` node `fn`. It uses the traversal described in Section 2.5.2 to iterate over the `Expr` nodes in the body. For each `Expr` node `expr`, the following actions are taken:

1. If `expr` is a `SymExpr` node, create a reference type for the type of the symbol if it doesn't yet exist.
2. Call `preFold` on `expr` to handle cases before we resolve further. This may change the node to which `expr` points. For example, if this is a call to the primitive `GET_REF`, then the compiler sees if the argument to the call is already a value. If it is, then the call is replaced by the argument and `expr` is updated to point to the argument. Numerous transformations are done during this step.
3. Resolve returned values of parameter functions.
4. Issue compiler errors when encountering the primitives wrapped by `compilerError` and `compilerWarning`.
5. If `expr` is a `CallExpr` node, update the resolution call stack and call `resolveCall` on `expr` (See Section 3.9.3.) If the call resolves to a function, then call `resolveFns` on that function.
6. Call `postFold` on `expr` to handle cases before we resolve further. This does many transformations on the AST and, like `preFold`, may change the node to which `expr` points, which will impact the traversal (the next node to which `expr` points).

3.9.3 `resolveCall(call)`

This function resolves the `CallExpr` node `call`. If the call is to an unresolved function, the algorithm proceeds as described in the language specification, more or less. If the call is to a primitive, transformations may be done on the AST similar to those done in `preFold`. In addition, the `MOVE` primitive is “resolved” in order to determine the type of any symbol on the left-hand side, or to do type checking if this symbol has already been resolved.

3.9.4 Wrappers

When resolving function calls, the actual arguments may not match the formal arguments in a one-to-one mapping with exact type matches. Such cases occur with implicit coercions, with default values for arguments, with named-argument passing, and with scalar promotion. To handle these cases, wrappers are created during function resolution, and marked with the “inline” flag. The wrappers are created using caches to ensure we don't create the same wrapper multiple times. This is necessary to avoid creating multiple instances of a type when there should just be a single type.

The default wrapper calls a function after inserting extra actual arguments into the call. The order wrapper calls a function after rearranging the arguments which can be in the wrong order due to named-argument passing. The coercion wrapper calls a function after casting arguments to other types. This includes dereferencing an actual argument that is a reference and reading the value in a sync or single type. The calls to `readFE` and `readFF` may be inserted in coercion wrappers. The promotion wrapper calls a function within a loop. The promotion wrapper is implemented as an iterator. Parallel promotion is handled by creating leader and follower iterators.

3.9.5 Generic Instantiation

Generic instantiation of functions occurs during function resolution. Calls to functions that are generic are resolved at the same time as calls that are not generic. When a generic call becomes a candidate for function resolution, it is instantiated and its where clause is evaluated. If it is selected, then it is resolved. This means, among other things, that calls in the where clause may be resolved even if the function is ultimately not resolved.

3.9.6 Runtime Types

Certain types in Chapel, such as array and domain types, are not purely static. They contain dynamic information. For example, the array type is largely static, being composed of an element type and a domain type, but the domain value is also part of the array type. This is even suggested by Chapel syntax where the value of the domain appears in an array type.

To handle these runtime types, the compiler replaces such types by values during function resolution. Thus if a function has a type argument and it is called with an array type, the type argument will not be eliminated from the function during function resolution but will rather be replaced by a value argument with a type that is the type of a runtime type value.

The pragma “has runtime type” is used to implement runtime types in a fairly general way. This pragma must be applied to the record type that has dynamic information associated with it and with a function that builds a value of that record type based on the arguments passed to the function. For example, for arrays, this is applied to a function as follows:

```
pragma "has runtime type"
def chpl__buildArrayRuntimeType(dom: domain, type eltType) type
    return dom.buildArray(eltType);
```

The runtime type is thus composed of two fields: a domain (value) and an element type (type). When creating a value from a runtime type, a call to this function will be inserted by the compiler, passing the values in the runtime type value to the arguments of this function.

In addition, a function called `chpl__convertValueToRuntimeType` must be defined that takes a value of the record and calls the function with the pragma. When determining the type type of a value of the record type, a call to this function will be inserted by the compiler.

In Chapel, in addition to arrays, domains have runtime types consisting of the value of their distribution. For sparse domains, the dense parent domain is also part of the runtime type.

The compiler does not currently correctly handle records or classes composed of fields that have runtime types. That is, if the type of a tuple of an array is passed to a function, the runtime type of the array component will be lost, to unfortunate effect.

3.9.7 Scalar Promotion Types

Scalar promotion keys off a type field of a class that identifies the type of scalar over which an aggregate promotes. During function resolution, this type is identified and then used when selecting candidates during overload resolution.

The implementation relies on a field called `_promotionType` in the Chapel type. For this reason, the implementation can easily be extended to user-defined types should the language be changed to expose such functionality to a user-defined class.

3.9.8 Conditional Resolution and Try Tokens

Function resolution can be done conditionally based on a “try token” called `gTryToken`. This global symbol can be used in the conditional expression of a conditional statement. When encountered by function resolution, the compiler will try to resolve the then-statement. If this is successful, the conditional statement will be replaced by the body of the then-statement. If not, the conditional statement will be replaced by the body of the else-statement, and resolution will continue on the else-statement as normal.

This is used to attempt to parallelize reductions. The parallel iterator is called in the then block. In the else block, the serial iterator is called and a `compilerWarning` call is made to indicate that the reduction has been serialized.

3.9.9 Iterator Handling During Resolution

The iterator is resolved as if it is a normal function where the yield statements are treated identically to return statements in terms of resolving the return type. In addition, the `iteratorInfo` field is created, as are stubs for the iterator class, the iterator record, and the methods and functions. These are described in more detail in Section 3.14, which also describes the `lowerIterators` pass.

A function `_getIterator` is created during function resolution that takes as an argument the iterator record and returns the iterator class. Such functions are resolved when actually iterating over a loop. In short, the reason to have both an iterator record and an iterator class is to make memory management simpler. It is easier to free the iterator classes since we just allocate them at the beginning of the loop. Iterator records are copied as necessary. The semantics of invoking an iterator multiple times after the iterator record has been created is ill-defined—this happens when an iterator is passed to a generic function and the captured argument is iterated over twice. When iterators are assigned to variables, an array is created. When iterators are passed to generic arguments, the iterator record is captured.

The links to the leader and follower iterators are created during function resolution by creating a function that calls the leader or follower from the `TO_LEADER` and `TO_FOLLOWER` primitives.

3.9.10 The Dream of Out-of-Order Resolution

There are a number of non-working tests (futures) that are expected to work should function resolution work if the strict ordering constraint on it now is removed. The dream is that the compiler could have a queue of functions that need to be resolved. Then if it fails for any reason, it could continue on the other functions. There are a number of places where odd calls to `resolveFns` or `resolveCall` could be removed.

This would be especially useful with code like:

```
module M1 {
  use M2;
  var a = 1;
  var b = c;
}

module M2 {
  use M1;
  var c = 2;
  var d = a;
}
```

In this example, the compiler is going to try and fail to resolve one of the module initialization functions first. If it could start resolving one, and then switch to the next when failing to resolve the type of either `b` or `d`, then it would be able to finish resolving the first later.

3.9.11 The Dream of Combined Scope and Function Resolution

In order to be able to declare variables or use modules in conditional statements where the conditional is a parameter such that the variable or module is visible in the outer scope, scope resolution and function resolution would have to be done at the same time. This would involve maintaining a symbol table while resolving function calls and types. It would also involve a bit of rearrangement of the compiler since there are a number of things that are done on the whole program between scope resolution and calls to `resolveFns` on particular functions.

3.10 checkResolved

This semantic checks pass flags errors if any of the following statements are true:

- Control may reach the end of a function that returns values.
- Functions return nested iterators (perhaps loop expressions).
- An enumeration constant is not a compile-time constant.

3.11 flattenFunctions

This pass collects all nested functions (the `DefExpr` of the `FnSymbol` is in another function, *i.e.*, the `parentSymbol` of the `defPoint` of the `FnSymbol` is another `FnSymbol`) into a vector of nested functions that is then passed to the following function:

```
void flattenNestedFunctions(Vec<FnSymbol*>& nestedFunctions);
```

This function denests all of the functions in the vector argument. Outer variables are identified and passed to the nested functions by reference. This is done in total because one nested function could call another nested function that needs additional outer variable references.

This functionality is encapsulated so that it can be used elsewhere. In particular, this function is also called when inlining a loop body into a recursive iterator to flatten the nested function that implements the loop body during the `lowerIterators`

pass. Also, this function is also called during the `parallel` pass when nested functions are created to implement *begin*, *cobegin*, *coforall*, and *on* statements.

3.12 `cullOverReferences`

This pass has two distinct but related purposes. The first purpose is to replace all calls to the reference version of *var functions* that do not need the reference into calls to the value version of the same function. The value version can be arbitrarily different due to the existence of the implicit “setter” argument that is true if the var function is used as an lvalue and false otherwise. Therefore, it is essential that this transformation be correct—it is not merely an optimization.

The implementation of the code that completes this first purpose relies on def and use maps to let the compiler determine if a reference really needs to be a reference.

The second purpose is to remove all references of array wrapper records, domain wrapper records, and iterator records. This is essential since returning such a reference could result in a reference to something that is not on the stack. Meanwhile, we do not want to put these things on the heap since the classes within the wrapper records are already on the heap.

3.13 `callDestructors`

The primary purpose of this pass is to insert calls to destructors for values when they go out of scope, including records, arrays, and domains.

The first thing this pass does is to call `fixupDestructors`. This function inserts functionality into destructors. It inserts calls to the destructors of all value fields (because these should be called automatically). It also inserts calls to the destructor of the parent class.

The function `insertAutoDestroyTemps` inserts the calls to destructors for variables when control exits their declaration scope. This requires an analysis to determine if control may or must exit the scope.

A large amount of code handles functions that return records because the compiler wants to insert code to free the record when it goes out of scope but the record is going to return it to the callsite. This is handled by returning records through reference arguments. Then the compiler can free the record if it is assigned by value.

Lastly, this function builds up a function to call destructors on global variables.

3.14 `lowerIterators`

In summary terms, this pass lowers an iterator, be it a serial iterator or a follower (leaders are always inlined), into a class with methods where these methods can be called in a loop to implement the functionality of this iterator. Before each iteration, a “has more” method can be called to see if the iterator is finished. If there are more values to be iterated over, an “advance” method can be called to get to the next value and a “get value” method can be called to get the particular value. The class is used to save state. Both the “has more” and “get value” methods are very simple. The “advance” method looks like the original iterator function, but there is a jump table at the beginning to return to the points in the code that immediately follow yield statements, and fields in the class store state for repeated calls to this function.

For each resolved iterator in the Chapel program, including each instantiation of every generic iterator and each iterator that implements a particular promotion wrapper, a class called `IteratorInfo` is created. This class has the following fields:

- `IteratorTag tag` identifies whether this iterator is a serial iterator, a leader iterator, or a follower iterator.
- `FnSymbol* iterator` points to the original iterator function. During this pass, this original function is transformed into a function that returns an instance of the iterator record.
- `FnSymbol* getIterator` points to the `_getIterator` function created during function resolution that takes an iterator record and returns an iterator class.
- `ClassType* iclass` points to this iterator’s implementing class. This class is created during this pass. The iterator class is only instantiated when using the iterator in a loop. It is destroyed when the loop completes.

- `ClassType* irecord` points to this iterator's implementing record. This record is created during this pass. The record is constructed by the `getIterator` function. The distinction between the iterator record and the iterator class makes it easier to free memory optimally because only the iterator record is bandied about. The iterator class is used in a much more structured way.
- `FnSymbol* advance` is a method on the iterator class that updates the state of the iterator so that calls to the `getValue` method can return the next value.
- `FnSymbol* zip1`, `FnSymbol* zip2`, `FnSymbol* zip3`, and `FnSymbol* zip4` are methods that implement code specialized for iterators with a single yield in a single loop.
- `FnSymbol* hasMore` returns true if the iterator is not finished.
- `FnSymbol* getValue` returns the value that the iterator is currently ready to return.

If iterators are not zipped and are not recursive and have a single yield statement, we implement them with straightforward inlining. That is, inline the iterator in place of the loop, and replacing the yield and return statements with copies of the loop body.

3.14.1 Single Loop Iterator Optimization

The single loop iterator optimization is important for generating code with optimal control flow when a common class of simple iterators are zippered together in a loop, or when inlining is not done. This common class of simple iterators can be described as any iterator with a single yield statement that is immediately in a single loop statement that is immediately in the iterator function. Without this optimization, goto statements will be used with the control flow to implement the jump table. This optimization creates zip methods, labeled `zip1` through `zip4`, that implement the functionality of the sections of code marked in the following code:

```
def iterator() {
  // zip 1
  loop {
    // zip 2
    yield value;
    // zip 3
  }
  //zip 4
}
```

3.14.2 Recursive Iterators

Recursive iterators are identified. We used to avoid inlining recursive functions altogether, but the compiler now has initial capability to inline the loop body into the recursive function at the yield points. This could be especially clean with function pointers, but we currently inline the loop body, and move the iterator so that it is a nested function.

In making the recursive function call, care is taken to avoid inlining the function again at the place where the loop is invoked.

Recursive iterators with on-statements, perhaps a common case one day, require careful coding in order to avoid superfluous remote references. The arguments passed to the recursive call are copied locally.

3.14.3 Iterators and Local Blocks

Inlining iterators is a bit tricky in the presence of *local* blocks which impose restrictions to the code lexically with the local block. To ensure that the body of the loop is not restricted to local functionality, the compiler counts the number of local blocks around `yield` statements before inlining, and then inserts a corresponding number of *unlocal* blocks around the loop body.

When not inlining, the local blocks also need to be handled with care. To do this, the compiler fragments the local blocks within an iterator so that control flow constructs do not appear within the local blocks. That is, a single local block may be split into multiple local blocks to avoid having control flow within the local block.

3.14.4 Miscellaneous Notes

A few notes about the implementation:

- When iterators are being lowered, the compiler inserts coded, but bogus, `GET_MEMBER` primitives to access the fields of the iterator class. These bogus primitives use numbers instead of fields because the fields have not yet been created. This impacts functionality during this time since we cannot determine the type of such primitive expressions.
- The compiler uses live variable analysis to determine what local variables are live at the points where the iterator yields values. Only these variables need to be represented by fields in the iterator class. If live variable analysis is not used, the compiler creates a field for every local variable.

3.15 parallel

This pass originally transformed the code to interface with the tasking runtime (the first and last steps below), and that was all. From a historical perspective, this pass enabled multithreading (before its introduction, execution was serial), hence its name. That said, this pass has a terrible name.

The current implementation of this pass takes the following actions in the following order:

1. Transform all `begin`, `cobegin`, `coforall`, and `on` blocks into nested functions and flatten them. Prior to this pass, these constructs are represented as `BlockStmts` with a primitive `blockInfo` field.
2. Run the optimization called remote value forwarding on all flattened nested functions created above. This optimization looks for functions that take references and replaces them with values if possible. Flattening of nested functions introduces reference arguments to handle the outer variables (as described in Section 3.11). Thus the function created to handle an `on`-statement will take reference arguments to any variables declared outside of the `on`-statement.

This optimization is currently much more conservative than it should be. A reference argument is changed into a value argument if the reference is never written to and if there are no calls to any functions involving synchronization within the body of the function. These are flow-insensitive checks, but this optimization could benefit from being flow-sensitive. For example, it may be worthwhile to pass in both the reference and the value if it is read before it is written. In addition, if the reference is read before any call to a synchronizing functions, it can be passed as a value. Essentially what this optimization is doing is to move the read up to the point where the function is called (or where the `on`-statement is executed). This optimization uses `def` and `use` maps as well as the call graph.

3. Reprivatize privatized-object fields (array and domain descriptors may be privatized) in Iterator classes (`reprivatizeIterators`). That is, code is inserted when accessing such fields to get the local private copy. This is not necessary with arbitrary classes and records because their fields will never capture privatized objects. Instead, they will always capture privatized IDs which can then be changed into objects at the point they are used. The IDs map to the privatized copy of a class on any given locale. Iterator classes, on the other hand, are built up by the compiler. If the privatized object (as opposed to the ID) is added to the class, then we may be pointing at a remote object when there is a local object that we could be pointing at.
4. Move stack-allocated data onto the heap as necessary (`makeHeapAllocations`). This function uses `def` and `use` maps to trace references through `begin`- and `on`-statements (or rather the functions that implement them). The data that such references point to needs to be put on the heap. For `begin`-statements, this has to happen because the `begin`-statement can return before the function completes. For `on`-statements, this has to happen because remote communication can only involve data on the heap.

This function also takes care of inserting code to broadcast the values in global constants so that each locale can access these values directly. Other globals are put on the heap and set up so that the references on every locale other than locale 0 points to the heap on locale 0.

At the end of the function `makeHeapAllocations` there is a call to `freeHeapAllocatedVars` that duplicates some functionality in an effort to free such heap-allocated structures. This code is incomplete in a number of ways. A noticeable limitation is that we won't be able to free heap-allocated structures completely until the compiler inserts code to reference count the number of tasks that can access such structures. This should probably not be universal since there will be simple cases where the reference count will not be necessary and where the performance hit (time and space) should be avoided.

5. Handle the implementation of the `EndCount` class used to implement sync statements, including the implicit sync statement around `main`. Accesses to this class are implemented via primitives `GET_END_COUNT` and `SET_END_COUNT` up to this point. The compiler threads `EndCount` variables through functions by adding arguments to these functions.
6. Bundle arguments to functions that implement `begin`, `cobegin`, `coforall`, and `on` statements. The tasking runtime invokes such functions via a pointer to a function that expects one argument. Thus the compiler changes these functions so that they expect only one argument. Structs are created to capture the multiple arguments.

3.16 `prune`

The `prune` pass, run again after the pass called `localizeGlobals`, has a narrow focus. During this pass, the compiler identifies unused functions and types, and removes them from the AST.

3.17 `complex2record`

This pass replaces the primitive complex types `dtComplex` with records composed of two floating-point values, one for the real part and the other for the imaginary part. All occurrences of the primitive types are replaced with the new records. The primitives `GET_REAL` and `GET_IMAG` are replaced with the primitive `GET_MEMBER`. Henceforth in compilation, the primitive complex type is of no concern.

3.18 `removeUnnecessaryAutoCopyCalls`

This optimization pass removes redundant calls to the auto copy mechanism used to implement memory freeing on values, including arrays and domains. This pass attempts to match calls to the “auto copy” function with calls to the “auto destroy” function and cancel them out. Since the “auto copy” and “auto destroy” functions may increment and decrement sync variables for reference counting arrays and domains, this optimization can significantly improve performance.

3.19 `inlineFunctions`

This optimization pass inlines all functions that are marked by the “inline” flag. Line numbers are updated to the call site. Before a particular function is inlined, optimizations like copy propagation are called on that function to try to reduce the amount of code that has to be replicated.

This pass also calls the `collapseBlocks` and `removeUnnecessaryGotos` functions on every function. These functions are described in Section 2.6.10.

3.20 `scalarReplace`

This optimization pass replaces some variables of some record types by multiple variables, one for each field in the record. The order of replacement is based on a topological sort of the types to avoid replacing variables assigned to fields.

This pass calls `eliminateSingleAssignmentReference`, described in Section 2.6.10, to eliminate a case that would otherwise disable this optimization.

In addition to the fairly straightforward code needed to scalar replace records, this pass also tries to scalar replace classes that are used in a stylized way. In particular, this pass tries to scalar replace iterator classes. In general, scalar replacing classes would not work since they are references and not values. For iterators, this is an important step in enabling many of the more traditional optimizations run later.

3.21 `refPropagation`

This optimization pass calls `singleAssignmentRefPropagation`, as described in Section 2.6.10, on every function.

3.22 copyPropagation

This optimization pass calls `localCopyPropagation`, `deadVariableElimination`, and `globalCopyPropagation` in that order on every function. These functions are described in Section 2.6.10.

3.23 deadCodeElimination

This optimization pass calls `deadCodeElimination`, as described in Section 2.6.10, on every function.

3.24 removeWrapRecords

This optimization pass runs a variation of scalar replacement on the wrapper records for arrays and domains. Since these records only have one field, they can be removed completely and all uses of these types can be replaced by the field type. Unlike in scalar replacement where a type may have multiple fields, variables are replaced by only one other value by this optimization. The field in question is `_value`.

This is more involved because these wrapper records also contain a field called `_valueType`. However, this field is only used for type information during function resolution (since privatization changes the `_value` field into an integer). Therefore, before doing the replacement described above, references to `_valueType` are first eliminated. This code relies on dead code elimination, inlining, and copy propagation, so if any of these optimizations are disabled, this pass is skipped.

3.25 removeEmptyRecords

This optimization pass removes all empty record types, and all variables and arguments with these types.

3.26 localizeGlobals

This optimization pass creates temporaries at the top of functions to capture global constants. This optimization was especially important on the Cray XMT (TM). In any event, since constants in Chapel cannot be identified as constants in C, this seems like a worthwhile optimization. This optimization does not handle the case where a function that reads a global constant is called from another function in a critical loop.

This optimization is a good one for the argument against source-to-source compilation, or at least source-to-source compilation to C. If the compiler did handle the above case (a function called in a critical loop), it would have to pass the constant into the function as an argument. Thus all constants may be passed around to functions, but that is clearly not optimal either. Such thinking must be weighed against the advantages of source-to-source compilation, namely portability.

3.27 returnStarTuplesByRefArgs

This pass changes all functions that return star tuples into function that take, as arguments, references to these star tuples and assign the values into these references. This pass also changes all `SET_MEMBER`, `GET_MEMBER`, and `GET_MEMBER_VALUE` primitives into `PRIM_SET_SVEC_MEMBER`, `PRIM_GET_SVEC_MEMBER`, and `PRIM_GET_SVEC_MEMBER_VALUE` primitives. These are used on star tuples. The other primitives were used on star tuples which are stored as records with fields in the AST.

3.28 insertWideReferences

This pass introduces wide references and wide classes into the AST. A wide reference is a reference to something that may exist on a different locale. It is generated as a struct with two fields: a locale number and an address. In the AST, wide references are represented as records with two fields: a locale number and a reference. Since classes are references, they too can be wide. A wide class is represented as a record with two fields: a locale number and a class. In both cases, the address or class reference is only valid on the locale indicated by the locale number.

This pass consists of the following steps:

1. Create a wide reference type for every type and a wide class for every class. This includes both wide references to classes and wide references to wide classes. Fortunately, there are no references to references, references to wide references, wide references to references, or wide references to wide references. Because enough is enough!
2. Change all occurrences of references into wide references. The primitives do not have to change in substantial ways in that the primitives that work over references also work over wide references. There are some simplifications that are necessary, however. For example, the compiler inserts code to dereference wide references to wide classes in some primitives so that the double remote access will not happen in the same primitive.
3. Create a function to allocate all global variables on the heap on locale 0 and set up wide references to these global variables on other locales.
4. Transform the code in local blocks to check that wide references are local (unless checks are disabled). This optimization analyzes the code in local blocks for places where wide references or wide classes are accessed such that remote communication may occur. In this case, a check is inserted to ensure that the wide reference is local to the current locale, and a narrow reference is inserted to capture the address. Note that wide references can be moved into other wide references even if they are not local, since this does not require communication.
5. Call `narrowWideReferences` to replace wide references and wide classes with narrow references and narrow classes if the compiler can prove that this is legal to do.

3.29 optimizeOnClauses

This pass marks functions that implement on-statements with the “fast on” flag if they can be executed directly by the handler on the remote locale (rather than being handled by a separate thread on the remote locale). This requires that the code be simple (no synchronization, no remote memory accesses, etc.) and relatively fast.

3.30 insertLineNumbers

This pass inserts line numbers and filenames into functions and calls to these functions so that errors that show up in the internal modules will pinpoint code that the Chapel programmer should know something about. That is, if a primitive is called in an internal module, then this pass will add two arguments to the function containing the primitive call. If the call site to this function is in user code, that line number and filename will be added to the call. Otherwise, the compiler recurses to the function that that function is called in.

The compiler takes care to pass line number and filename information through argument bundles created during the `parallel` pass for functions called indirectly in the runtime.

3.31 repositionDefExpressions

This pass moves `DefExpr` nodes into the innermost `BlockStmt` nodes that they can legally be declared within. This decreases the scope of the declaration. This is essential for compiling with certain pragmas on the Cray XMT (TM) where declarations must be inside a parallel loop to avoid races due to assigning values to variables shared between iterations.

This optimization currently looks at all occurrences of local variables in a function and determines what block these variables should be declared within based on these occurrences. This may be insufficient if the compiler ever generates code where a variable may be live in a block outside of any of its uses. With the iterator transforms, and its loops and goto statements, this may indeed be possible.

3.32 codegen

This pass generates C code from the AST. All declarations, function prototypes, etc., are placed in a single header file called `chpl__header.h`. C code for each module is put into a separate C file. All of these files are included by a file called `_main.c` which is compiled in the `makeBinary` pass.

To compile each module, the compiler invokes the `codegenDef` method on the `ModuleSymbol` node. The recursive functions `codegenDef` and `codegen` generate the code. This design makes fewer assumptions about the normalized form than it can. In particular, it assumes there may be deeper recursion than will actually occur given normalization.

One of the more elaborate parts of code generation relates to the generation of primitives. This is complicated by the reference semantics since the code we generate for many primitives depends on whether the type of the argument is a reference or not.

3.33 `makeBinary`

This pass invokes the C compiler and linker via the Makefile created during the `codegen` pass.

3.34 What Happens Between Passes?

There are two main actions that take place between every pass. First, the compiler traverses the AST and the global vectors of all AST nodes, and removes all AST nodes that are not in the AST, reclaiming the memory. Second, the compiler verifies, via `verify` methods on all node types, that the AST is in a coherent state. This could be turned off for non-developers to save time.

In addition, each pass is timed and these times are printed if using the `-print-passes` flag and statistics are gathered and printed if using the `-print-statistics` flag.

4 Miscellaneous

4.1 Compiler Strings

Strings are canonicalized via the function `astr`. This function can take up to 8 string arguments that are concatenated together. The strings are stored in a large hash table called `chapelStringsTable`. During compilation, we typically canonicalize strings which allows them to be compared via a pointer comparison. The hash table of strings is freed when the compiler completes.

Additionally, the function `istr` can convert an integer to a canonicalized string.